

DaCaF: Bayes-optimal per-metric inference for probabilistic multi-label classification

Vu-Linh Nguyen^a, Xuan-Truong Hoang^{b,*}, Anh Hoang^c, Van-Nam Huynh^b

^a *TODO: organization, address, country*

^b *Japan Advanced Institute of Science and Technology (JAIST), Nomi, Ishikawa, Japan*

^c *TODO: organization, address, country*

* Corresponding author. E-mail: hxtruong@jaist.ac.jp

Abstract

In multi-label classification (MLC), different evaluation metrics call for different Bayes-Optimal Predictions (BOPs), and no single model is optimal for all of them. Probabilistic MLC allows the target metric to be optimized during the prediction phase, but computing the exact BOP is intractable in general, and existing tools are limited—covering only a few metrics, relying on strong assumptions such as label independence, or lacking a reusable, well-tested implementation. We present DaCaF, a Python package implementing the Divide-and-Conquer and Fusion approach that, given any probabilistic model $P(y | x)$, returns the exact BOP for a chosen metric, covering two families of MLC metrics. The package provides seven per-metric prediction rules, each verified against brute-force enumeration, together with a command-line interface, a library API, and reproducible evaluation scripts.

Keywords

Multi-label classification; Probabilistic classifier chains; Bayes-optimal inference; Evaluation metrics

Code metadata

Nr.	Code metadata description	Please fill in this column
C1	Current code version	v1.1.0
C2	Permanent link to code/repository used for this code version	https://github.com/hxtruong6/inference_probabilistic_mlc Zenodo: https://doi.org/10.5281/zenodo.20572638 PyPI: https://pypi.org/project/dacaf-mlc/
C3	Permanent link to Reproducible Capsule	https://codeocean.com/capsule/1580907/tree
C4	Legal Code License	MIT License
C5	Code versioning system used	git
C6	Software code languages, tools, and services used	Python (≥ 3.10); NumPy, SciPy, scikit-learn, pandas, joblib; optional PyTorch / torchvision / torchxrayvision ([image] extra)
C7	Compilation requirements, operating environments & dependencies	<code>pip install dacaf-mlc</code> (core) or <code>"dacaf-mlc[image]"</code> ; Linux/macOS; Docker image provided
C8	If available Link to developer documentation/manual	README.md, CONVENTIONS.md, docs/
C9	Support email for questions	hxtruong@jaist.ac.jp

1. Motivation and significance

Multi-label classification (MLC) is a setting in which each instance can belong to any subset of the labels, with applications ranging from text and image annotation to high-stakes domains such as clinical decision-making. To compare ground-truth label sets and predictions, a variety of evaluation metrics with different characteristics have been proposed. Probabilistic MLC offers unique advantages by allowing the optimization of different metrics during the prediction phase: once a probabilistic classifier estimates the conditional probability of each labeling, the Bayes-Optimal Prediction (BOP) of a target metric—the prediction $\hat{y}$ that maximizes the expected score of that metric—can be sought. As different metrics may call for significantly different BOPs, a mismatch between the target metric, i.e., the metric optimized during the prediction phase, and the evaluation metric, i.e., the metric used to assess the classifier, can negatively impact predictive performance.

However, computing the BOP in a naive way—by searching over the 2^L candidate predictions for L labels—quickly becomes intractable as the number of labels grows. Existing studies are arguably limited in two respects: some focus on only a small number of metrics, such as subset 0/1 accuracy, Hamming accuracy, and F_β (Dembczyński et al., 2010, 2011; Waegeman et al., 2014), yet require a different, metric-specific algorithm for each metric; others rely on strong assumptions, such as label independence (Pillai et al., 2017). As a consequence, no reusable, exact, and broadly applicable tool exists for finding BOPs across the metrics commonly used in probabilistic MLC.

We present DaCaF, a software package that implements the Divide-and-Conquer and Fusion (DaCaF) approach (Nguyen et al., 2026): given any probabilistic model $P(y | x)$, it returns the exact BOP for a chosen metric, covering two families of commonly used MLC metrics. It proceeds in two steps—a divide-and-conquer strategy followed by a fusion strategy—to find the prediction that is optimal for the selected metric. DaCaF thereby (a) optimizes the metric one actually cares about, (b) enables studying the mismatch between the target and evaluation metrics in an exact way, undisturbed by approximation, and (c) provides a diagnostic that flags metrics whose BOP has a trivial or very restricted structure—useful for assessing the practical value of emerging MLC metrics.

2. Software description

2.1 Software architecture

DaCaF is an installable Python package (`dacaf_mlc`; pip / Docker, with a console script `dacaf-mlc`). Its core class `ProbabilisticClassifierChainCustom` (PCC) wraps any scikit-learn base estimator and estimates $P(y | x)$ via a classifier chain. Every inference rule rests on two algorithmic building blocks:

- *Divide-and-conquer*: partition the 2^L predictions into $L+1$ groups by number of relevant labels; the best in each group is found by sorting labels by a per-metric score, and the global best is the best across groups.
- *Fusion*: the scores require marginal/pairwise probabilities, estimated by fusing the chain’s dependent binary classifiers via ancestral sampling.

Supporting modules include `evaluation_metrics.py`, `metrics_registry.py`, `pipeline.py` (train / k-fold / run_single), `evaluate.py` (CLI), `arff_dataset.py` (MULAN ARFF + 10-fold CV), `datasets.py`, and `chest_xray_dataset/` (NIH feature extractor, [image] extra). A vendored `skmultiflow` `ClassifierChain` base keeps the install self-contained.

2.2 Software functionalities

- Seven per-metric Bayes-optimal rules: `predict_fmeasure`, `predict_hamming`, `predict_markedness`, `predict_precision`, `predict_npv`, `predict_recall`, `predict_subset`. Each returns the prediction maximising the expected value of its target metric (conventions in `CONVENTIONS.md`).
- Covers two whole metric families at once and flags when a metric’s optimum is *trivial* (a metric-choice diagnostic).
- Batched predictor: one `predict_proba` per chain level instead of $N \cdot L \cdot 2^L$; verified numerically equal to brute-force enumeration.
- Reproducible evaluation: a CLI and scripts (`reproduce_tabular.sh`, Slurm) over MULAN tabular datasets and an NIH ChestX-ray14 image dataset.
- Correctness harness: every rule checked against brute-force expected-metric enumeration on small cases; CI on Python 3.10/3.12.

2.3 Illustrative example

A minimal library usage fits the model once and then queries the BOP for any metric:

```
from sklearn.linear_model import LogisticRegression
from dacaf_mlc.probability_classifier_chains \
    import ProbabilisticClassifierChainCustom
from dacaf_mlc.evaluation_metrics import EvaluationMetrics as EM

pcc = ProbabilisticClassifierChainCustom(
    LogisticRegression(max_iter=10_000))
pcc.fit(X_train, Y_train)           # Y: (n, L) binary
y_f1 = pcc.predict_fmeasure(X_test, beta=1) # BOP for F1
y_ham = pcc.predict_hamming(X_test)        # BOP for Hamming
print(EM.f_beta(Y_test, y_f1), EM.hamming(Y_test, y_ham))
```

The headline experiment (Table 1) shows that the mismatch between the target and evaluation metrics hurts: rows are the metric optimised for, columns the metric evaluated, and the bold diagonal entry is the maximum of its column. A single command reproduces it: `make reproduce` (CHD-49, the fastest dataset). The chest-xray path scales the same API to real DenseNet/ResNet features.

Table 1: CHD-49, PCC + logistic regression, mean over 5 seeds $\times$ 10-fold CV (% , higher is better). Read each column: the bold diagonal entry (optimise the metric you evaluate) is the maximum of its column. NPV and Recall collapse to the predict-all trivial optimum.

Target $\downarrow$ / Eval $\rightarrow$	F_1	Hamming	Markedness	NPV	Precision	Recall	Subset
Predict F_1	67.07	69.29	71.91	81.09	62.74	79.22	15.06
Predict Hamming	63.92	70.78	71.27	75.74	66.62	66.93	18.05
Predict Markedness	34.20	63.70	76.96	71.10	33.37	40.45	8.39
Predict NPV	58.35	43.01	71.50	100.00	43.01	99.46	0.00
Predict Precision	40.54	64.83	68.31	63.10	73.52	29.10	3.10
Predict Recall	58.35	43.01	71.50	100.00	43.01	99.46	0.00
Predict Subset	63.96	69.37	70.35	76.47	64.15	69.79	18.92

3. Impact

DaCaF enables an exact research question that approximation would otherwise blur: how much does optimising the wrong metric cost? Because the package is model-agnostic—it works with *any* probabilistic MLC model as a drop-in BOP layer on top of existing $P(y | x)$ estimators—this question can be studied across estimators and datasets without changing the inference code. Its breadth is unusual: a single tool covers two whole metric families, and the trivial-optimum diagnostic guides metric selection by flagging metrics whose Bayes-optimal prediction collapses to a degenerate form (as NPV and Recall do in Table 1).

The software is designed for reuse beyond the companion paper. Its closed-form rules are verified against brute-force enumeration and kept reproducible by CI, so results can be trusted and rebuilt. The barrier to entry is low: the package is pip-installable and Dockerised, ships scripted reproduction and bundled MULAN datasets, and lets a newcomer reproduce a paper table in minutes. A registry pattern makes it straightforward to add metrics and datasets, while the vendored base estimator keeps installs self-contained. Beneficiaries include MLC researchers studying metric behaviour, applied ML practitioners choosing a deployment metric, and educators teaching Bayes-optimal decision-making in MLC; an NIH ChestX-ray14 demonstration shows the same API applied to clinical/image MLC.

4. Conclusions

DaCaF turns a proved theoretical recipe into a reusable, verified, and exact per-metric inference tool for probabilistic MLC. Future work includes supporting more base models and metric families, and adding optional approximate fast paths for very large label spaces.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

TODO statement (or remove if not applicable). Template wording: “During the preparation of this work the author(s) used [TOOL] in order to [REASON]. After using this tool/service, the author(s) reviewed and edited the content as needed and take(s) full responsibility for the content of the published article.”

Acknowledgements

TODO funding sources (use the Guide’s standard wording), or state none.

References

References

- V.-L. Nguyen, X.-T. Hoang, A. Hoang, V.-N. Huynh, Probabilistic multi-label classification via a divide-and-conquer and fusion approach, *Information Fusion* (2026), Article 104517. <https://doi.org/10.1016/j.inffus.2026.104517>.
- X.-T. Hoang, V.-L. Nguyen, A. Hoang, V.-N. Huynh, DaCaF: probabilistic multi-label classification (v1.1.0) [Computer software], Zenodo (2026). <https://doi.org/10.5281/zenodo.20572638>.
- K. Dembczyński, W. Cheng, E. Hüllermeier, Bayes optimal multilabel classification via probabilistic classifier chains, *ICML* 2010.
- K. Dembczyński, W. Waegeman, W. Cheng, E. Hüllermeier, An exact algorithm for F-measure maximization, *NeurIPS* 2011.
- W. Waegeman, K. Dembczyński, A. Jachnik, W. Cheng, E. Hüllermeier, On the Bayes-optimality of F-measure maximizers, *J. Mach. Learn. Res.* 15 (2014) 3513–3568.
- I. Pillai, G. Fumera, F. Roli, Designing multi-label classifiers that maximize F measures: state of the art, *Pattern Recognition* 61 (2017) 394–404.